

Benchmarks of **redis-py**, **hiredis-py** and string creation

@mokaddem

March 5, 2018

Contents

1	Aggregated results	3
1.1	Summary	3
1.1.1	Pushing - LPUSH	3
1.1.2	Popping - RPOP or LRange/LTrim	3
2	Detailed implementations and tests	5
2.1	Benchmark parameters	5
2.2	Naive implementation	6
2.3	Pipeline feature	7
2.4	Replacing RPOP by LRange/LTrim	8
2.5	Using <code>redis-cli --pipe</code>	9
2.6	Generating redis protocol	10
2.6.1	Generating redis protocol with <code>string.Formatter</code>	10
2.6.2	Generating redis protocol with <code>String</code> + concatenation operator	10
2.6.3	Generating redis protocol with <code>String</code> % substitution operator	11

Introduction

In the course of a software development, I needed a buffering medium which was able to fulfil these requirements:

- Capable of handling millions of request per second
- Data persistence
- Easy to use

Where Redis suits these.

My software being written in python, I evidently used one of the redis python client available; **redis-py** is the recommended one.

This report focuses only on throughput, so data size was fixed and was fitting my needs.

It contains benchmarks of **redis-py**, where we can see that it is particularly slow compared to the c-written **hiredis** library or **redis-benchmark**. It also presents some features and technics to boost performances, as well as comparing methods to rapidly create strings.

Chapter 1

Aggregated results

1.1 Summary

These tables and plots show a summary for each tested different configuration.

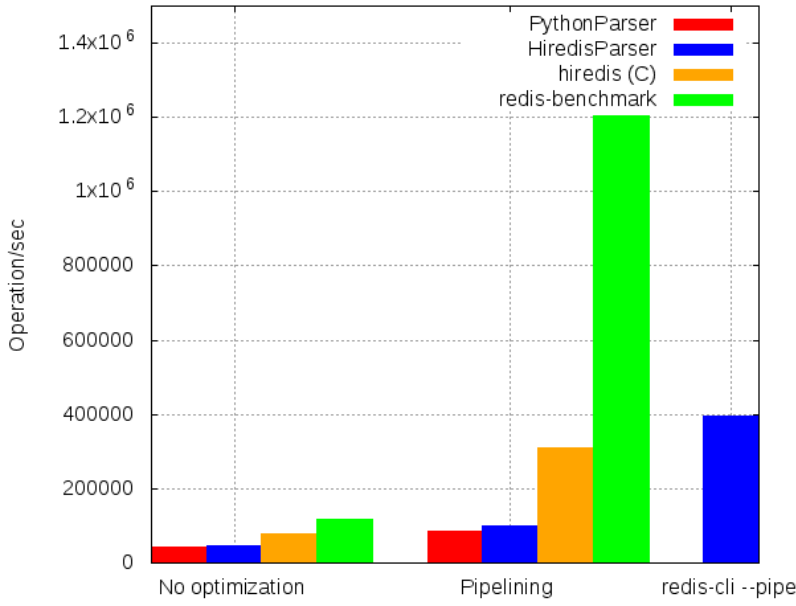
1.1.1 Pushing - LPUSH

	Tool/feature	Parser	Total time taken (s)			Operation/sec		
			worst	best	average	worst	best	average
2.2	Naive	PythonParser	2.86	2.20	2.37	34,942	45,464	42,158
	Naive	HiredisParser	2.20	2.09	2.17	45,388	47,745	46,150
	Naive	hiredis (C) ¹	1.33	1.23	1.25	75,363	81,365	79,977
2.3	Pipelining -P 1000	PythonParser	1.18	1.16	1.17	85,078	85,979	85,642
	Pipelining -P 1000	HiredisParser	1.01	0.96	0.98	98,863	104,368	101,535
	Pipelining -P 1000	hiredis (C)	0.35	0.31	0.32	281,466	320,858	308,386
2.5, 2.6.1	redis-cli, format	HiredisParser	0.28	0.23	0.25	358,529	437,094	394,046
2.6.2	redis-cli, concatenation +	HiredisParser	0.25	0.21	0.22	400,002	487,566	454,178
2.6.3	redis-cli, substitution %	HiredisParser	0.16	0.15	0.16	610,237	687,978	638,709
-	redis-cli, (no generation)	HiredisParser	0.08	0.07	0.07	1,287,382	1,523,811	1,400,926
	redis-benchmark -c 1 -d 129		-			-		
	redis-benchmark -c 1 -d 129 -P 1000		-			-		

1.1.2 Popping - RPOP or LRANGE/LTRIM

	Tool/feature	parser	Total time taken (s)			Operation/sec		
			worst	best	average	worst	best	average
2.2	Naive	PythonParser	2.68	2.11	2.27	37,301	47,323	44,021
	Naive	HiredisParser	1.96	1.85	1.91	51,050	53,950	52,383
	Naive	hiredis (C)	1.10	1.05	1.07	90,656	95,411	93,724
2.3	Pipelining	PythonParser	1.10	1.09	1.09	90,862	91,730	91,402
	Pipelining	HiredisParser	0.83	0.77	0.80	121,207	129,375	124,924
	Pipelining	hiredis (C)	0.15	0.13	0.14	659,078	751,359	720,310
2.4	LRANGE trick	PythonParser	0.34	0.32	0.33	292,466	311,161	304,744
	LRANGE trick	HiredisParser	0.08	0.07	0.07	1,207,644	1,511,085	1,416,695
	LRANGE trick	hiredis (C)	0.05	0.04	0.04	2,152,203	2,427,007	2,318,716
	redis-benchmark -c 1 -d 129 (RPOP)		-			-		
	redis-benchmark -c 1 -d 129 -P 1000 (RPOP)		-			-		
	redis-benchmark -c 1 -d 129 (LRANGE.100)		-			-		

¹Official C client




LPUSH:

With no optimization, python performs close to 60% of hiredis (C) performance:

$$\frac{\text{python perf}}{\text{hiredis (C) perf}} = \frac{46,150}{79,977} = 57.70\%$$

With pipelining, python performs close to only 30% of hiredis (C) performance:

$$\frac{\text{python perf}}{\text{hiredis (C) perf}} = \frac{101,535}{308,386} = 30.91\%$$



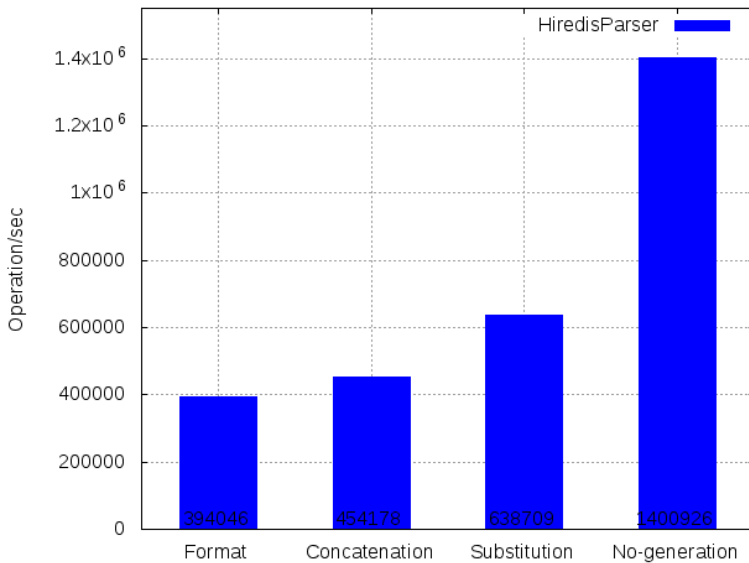
POP:

With no optimization, python performs close to 60% of hiredis (C) performance::

$$\frac{\text{python perf}}{\text{hiredis (C) perf}} = \frac{52,383}{93,724} = 55.89\%$$

With pipelining, python performs close to only 13% of hiredis (C) performance:

$$\frac{\text{python perf}}{\text{hiredis (C) perf}} = \frac{91,402}{720,310} = 12.69\%$$



String Generation:

Method	Cummulative gain
format	0%
Concatenation	15.26%
Substitution	40.63%

Chapter 2

Detailed implentations and tests

2.1 Benchmark parameters

All benchmarks have been performed using these parameters

- Usage of redis `unix_socket_path`
- Payload: size = 129 bytes

```
1 '{"origin": null, "channel": 0, "content": "redis@tshark_save:53619abd-a27c-432c-8f8d-1d059aab5f24", "size": 54, "redirect": true}'
```

- Operations:
 - 100000 LPUSH
 - Number of POP may differ if we use RPOP or LRANGE/LTRIM
- For each tests, we are using two different parsers for responses: PythonParser and HiredisParser¹
- In order to simulate a working software popping and adding elements to a buffer, the logic of the benchmark is the folloing:

```
1 # push
2 # c = total count      = 100,000
3 # d = divisor          = 1,000
4 # c/d = iteration count = 100
5 for i in range(int(c/d)):
6     for j in range(d):
7         # push
8     for j in range(d):
9         # pop
```

- For each benchmark, the processing has been done 10 times, then averaged
- For each 100 $\left(\frac{100,000}{1,000}\right)$ iterations, we are pushing and popping 1,000 elements

¹Hiredis is a C library that is available with Python bindings, `redis-py` will attempt to use the HiredisParser if you have the hiredis module installed and will fallback to the PythonParser otherwise.

2.2 Naive implementation

For each payloads to be buffered, we push them immediatly. Then, we retrieve them one at a time with a simple POP command.

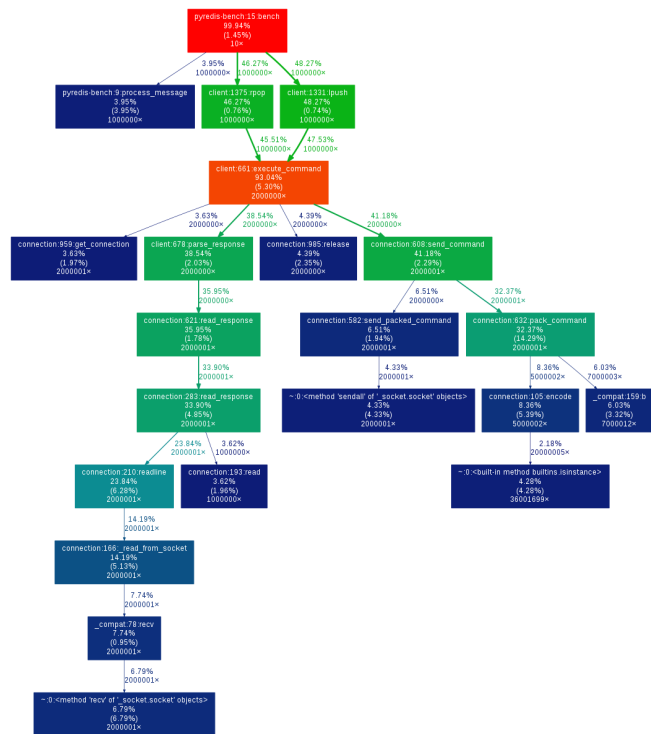
Source code

```
1 # push
2 t1=time.time()
3 for j in range(d):
4     redis.lpush('k', payload)
5     cpush+=1
6 time_push += time.time()-t1
7 # pop
8 t1=time.time()
9 for j in range(d):
10    msg = redis.rpop('k')
11    process_message(msg)
12    cpop+=1
13 time_pop += time.time()-t1
```

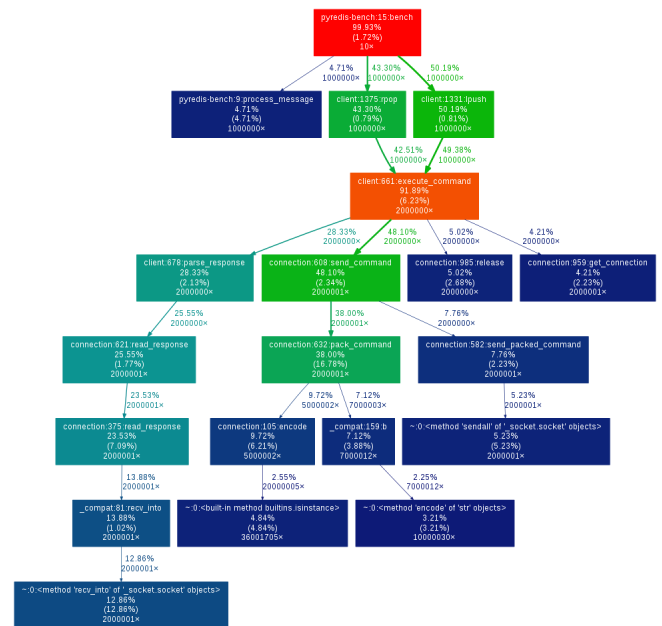
Profiling summary

Cumulative CPU %	PythonParser	HiredisParser
Processing	3.95%	4.71%
Receiving response	38.54%	28.33%
Sending command	41.18%	48.10%
Parser total gain: 19.24%		

PythonParser



HiredisParser



2.3 Pipeline feature

Here, we use the redis pipeline² feature. For each payload to be buffered, we push them in a pipeline which will execute pending command every 1000 pushing operation. Then, we retrieve them by sending the POP command in a pipeline which will also execute pending command every 1000 popping operation.

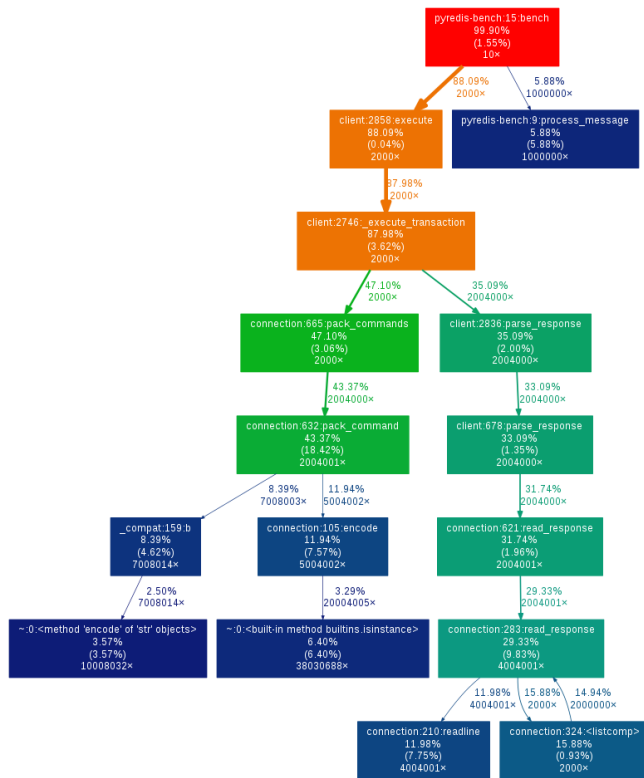
Source code

```
1 # push
2 t1=time.time()
3 for j in range(d):
4     pipeline.lpush('k', payload)
5     cpush+=1
6 pipeline.execute()
7 tpush += time.time()-t1
8 # pop
9 t1=time.time()
10 for j in range(d):
11     pipeline.rpop('k')
12 resp = pipeline.execute()
13 for msg in resp:
14     process_message(msg)
15     cpop+=1
16 tpop += time.time()-t1
```

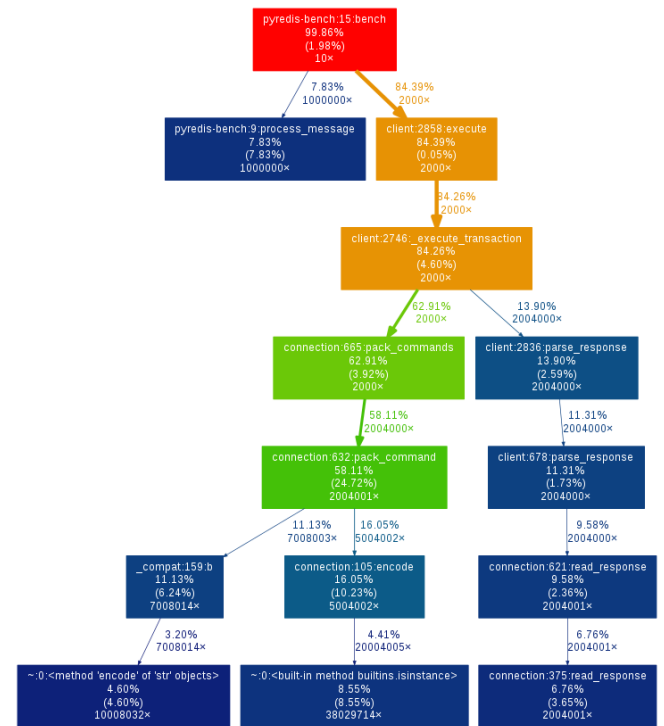
Profiling summary

Cumulative CPU %	PythonParser	HiredisParser
Processing	5.88%	7.83%
Receiving response	35.09%	13.90%
Sending command	47.10%	62.91%
Parser total gain:	33.16%	
Gain compare to 2.2:	48.86%	66.24%

PythonParser



HiredisParser



²Pipelining allows to send multiple commands to the server without waiting for the replies, and finally read the replies in a single step

2.4 Replacing RPOP by LRANGE/LTRIM

In this implementation, we are still using the redis pipeline feature for pushing data into the buffer, but we modify the popping behavior. Instead of sending one POP command at a time to the pipeline, we fetch a range (100 items) of buffered data with LRANGE, then we trim the buffer so that it mirror the effect of a POP command.

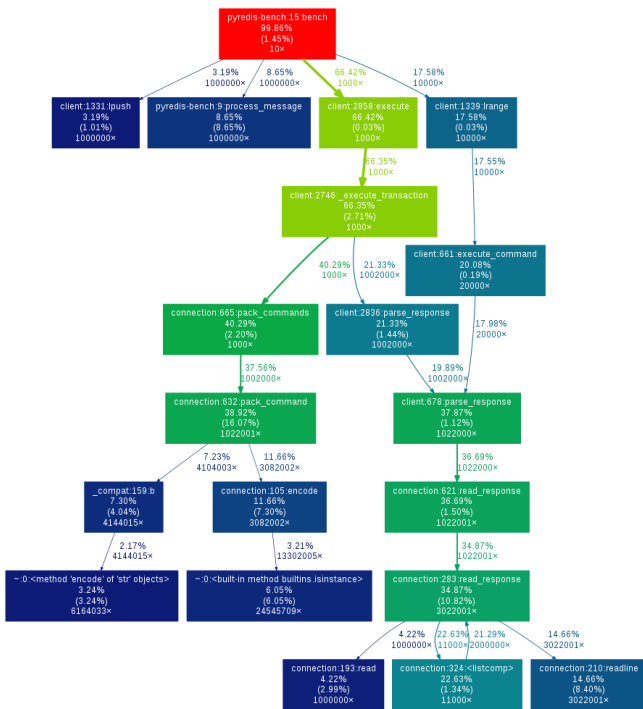
Source code

```
1 # lrange_count = 100
2 # push
3 t1=time.time()
4 for j in range(d):
5     pipeline.lpush('k', payload)
6     cpush+=1
7 pipeline.execute()
8 tpush += time.time()-t1
9 # pop
10 t1=time.time()
11 for j in range(int(d/lrange_count)+1):
12     msg_list = redis.lrange('k', -
13                             lrange_count, -1)
14     redis.ltrim('k', 0, -len(msg_list)-1)
15     for msg in msg_list:
16         process_message(msg)
17     cpop+=1
18 tpop += time.time()-t1
```

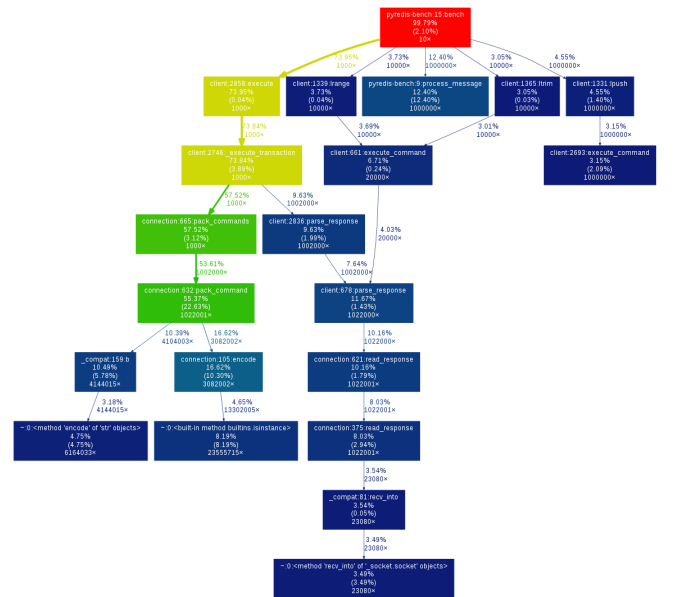
Profiling summary

Cumulative CPU %	PythonParser	HiredisParser
Processing	8.65%	12.40%
Receiving response	37.87%	11.67%
Sending command	40.29%	57.52%
Parser total gain:	43.35%	
Gain compare to 2.3:	47.11%	58.37%

PythonParser



HiredisParser



2.5 Using redis-cli --pipe

In this implementation, we are still using the redis pipeline feature and the LRange/LTRIM replacement of POP. However, we will use the `redis-cli` binary (provided by redis) with the `--pipe` options³.

In order to use this binary, we have to generate the valid redis protocol from the wanted command. which will be written to the STDIN of `redis-cli`.

Source code

```
1 # push with redis-cli
2 t1=time.time()
3 for j in range(d):
4     write_to_stdin(generate_redis_protocol(
5         'lpush', 'k', payload))
6     cpush+=1
7 tpush += time.time()-t1
8 # pop with pipeline
9 t1=time.time()
10 for j in range(int(d/lrange_count)+1):
11     msg_list = redis.lrange('k', -
12         lrange_count, -1)
13     for msg in msg_list:
14         process_message(msg)
15     cpop+=1
16 tpop += time.time()-t1
17 # flush and close stdin
18 # pop newly flushed items
```

Profiling summary

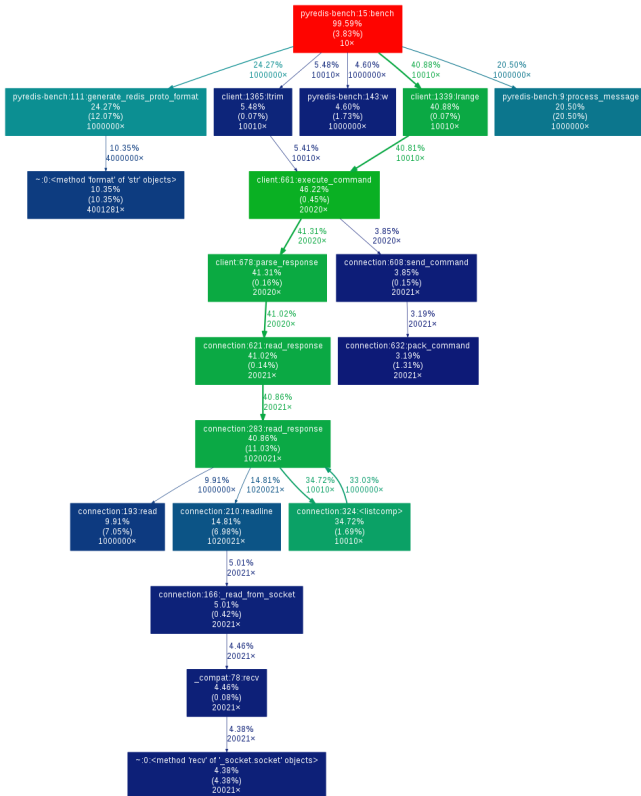
Cumulative CPU %	PythonParser	HiredisParser
Processing	20.50%	37.08%
Receiving response	41.51%	6.96%
Sending command ^a	8.45% ^b	14.79% ^c
Parser total gain:	80.87%	
Gain compare to 2.4:	136.99%	199.03%

^apopping + mass insertion with `redis-cli`

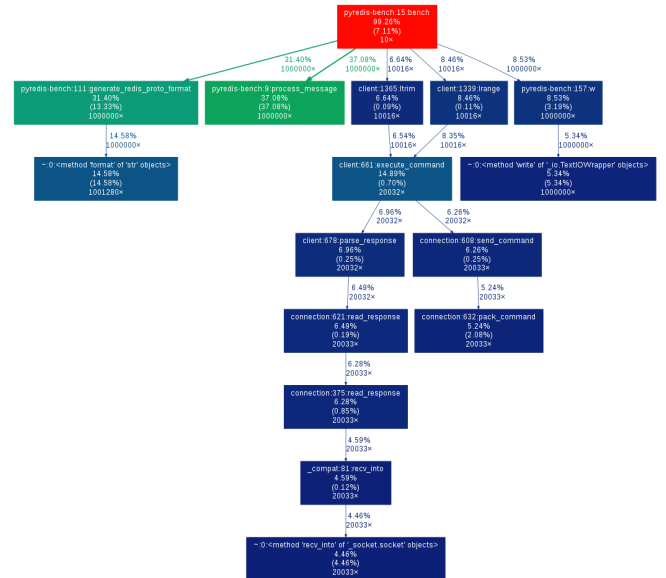
^b3.85% + 4.6%

^c6.26% + 8.53%

PythonParser



HiredisParser



³`redis-cli --pipe` consists in writing new commands while you read replies at the same time, ignoring the round trip time for every command

2.6 Generating redis protocol

We saw in section 2.5 that we need to generate the redis protocol ourself. We will now see different implementations of doing so, each one with drastic performance improvement.

For each algorithm, we use the same source code as in 2.5 with the `HiredisParser`.

2.6.1 Generating redis protocol with `string.Formatter`

```
1 def generate_redis_proto_format(cmd, key, value=''):
2     cmd_split = cmd.split()
3     if value == '':
4         proto = '*{argNum}\r\n${argLen1}\r\n{arg1}\r\n${argLen2}\r\n{arg2}\r\n'.format(
5             argNum=3 if value != '' else 2,
6             argLen1=len(cmd), arg1=cmd,
7             argLen2=len(key), arg2=key
8         )
9     else:
10        proto = '*{argNum}\r\n${argLen1}\r\n{arg1}\r\n${argLen2}\r\n{arg2}\r\n${argLen3}\r\n{
11        arg3}\r\n'.format(
12            argNum=3 if value != '' else 2,
13            argLen1=len(cmd), arg1=cmd,
14            argLen2=len(key), arg2=key,
15            argLen3=len(value), arg3=value)
16    return proto
```

Profiling summary

It should be noted that the cumulative CPU time of `process_message` entirely depends on the function implementation. Still, it is given to starting point for comparison.

Cumulative CPU %	HiredisParser
<code>process_message</code>	37.08%
<code>generate_redis_protocol_format</code>	31.40%
Generation loss ⁴	46.87%

2.6.2 Generating redis protocol with `String + concatenation operator`

```
1 def generate_redis_proto_concat(cmd, key, value=''):
2     cmd_split = cmd.split()
3     proto = '*' + (str(3) if value != '' else str(2)) + '\r\n'
4     proto += '$' + str(len(cmd)) + '\r\n' + cmd + '\r\n'
5     proto += '$' + str(len(key)) + '\r\n' + key + '\r\n'
6     if value != '':
7         proto += '$' + str(len(value)) + '\r\n' + value + '\r\n'
8     return proto
```

Profiling summary

Cumulative CPU %	HiredisParser
<code>processing</code>	40.53%
<code>generate_redis_protocol_string</code>	26.34%
Gain compare to 2.6.1:	9.30%
Generation loss	34.26%

⁴Cost of generating the protocol compared to not generating it at all.

2.6.3 Generating redis protocol with String % substitution operator

```
1 def generate_redis_proto_subst(cmd, key, value=''):
2     cmd_split = cmd.split()
3
4     if value != '':
5         proto = '%s\r\n%s\r\n%s\r\n%s\r\n%s\r\n%s\r\n%s\r\n%s\r\n' % ((str(3) if value != '' else
6             str(2)), len(cmd), cmd, len(key), key, len(value), value)
7     else:
8         proto = '%s\r\n%s\r\n%s\r\n%s\r\n%s\r\n' % ((str(3) if value != '' else str(2)), len
9             (cmd), cmd, len(key), key)
10    return proto
```

Profiling summary

Cumulative CPU %	HiredisParser
processing	43.62%
generate_redis_protocol_subst	19.92%
Gain compare to 2.6.2:	7.62%
Generation loss	24.85%

Conclusion

We saw that it is not recommended to use **redis-py** if we need pure throughput. Even so, some features like pipelining or using a different parser increase reasonably the performance, we are far from those achieved by **hiredis** (C).

In anycase, using the **HiredisParser** increased performances up to 40% for our case. Don't forget: **redis-py** uses **HiredisParser** if it is installed, therefore:

```
1 pip install redis
2 pip install hiredis
```